

A Limited Memory Steepest Descent Method¹

Yan Cheng²

April 19, 2021

¹Fletcher, R. A limited memory steepest descent method. Math. Program. 135, 413-436 (2012). <https://doi.org/10.1007/s10107-011-0479-6>

²Department of Applied Physics and Applied Mathematics, Columbia University

A Brief Review

From class, we know the steepest descent method computes

$$p_k = -\nabla f(x_k)$$

and uses linear search to update

$$x_{k+1} = x_k + \alpha_k p_k.$$

The SDM has certain advantages. However, it is known to be inefficient.

Motivation

NLP \implies SLP
Large-scale Nonlinear Programming \implies Sequential Linear Programming

Some reasonable requirements:

- ▶ Not have to store Hessian
- ▶ Can somehow make use of information from previous SLP iterations

Some possibilities:

- ▶ Conjugate-Gradient method
- ▶ 1-BFGS

Some possibilities:

- ▶ Conjugate-Gradient method
- ▶ 1-BFGS
- ▶ Or could we perhaps reconsider the steepest descent method?

Some step-length choices for SDM

If A is the Hessian of f , then

► (Line Search)

$$\alpha_k = \frac{p_k^T p_k}{p_k^T A p_k}$$

► (Barzilai-Borwein)

$$\alpha_k = \frac{p_{k-1}^T p_{k-1}}{p_{k-1}^T A p_{k-1}}$$

Here $p_k = -\nabla f(x_k)$ is the search direction in the k -th iteration.

Quadratic Problem

► Problem:

$$\min f = \frac{1}{2} \mathbf{x}^T A \mathbf{x}$$

where A is symmetric, positive definite. For the rest of this talk, denote

$$g_k = \nabla f(x_k).$$

- Using the steepest-descent update formula

$$x_{k+1} = x_k - \alpha_k g_k \iff x_{k+1} - x_k = -\alpha_k p_k,$$

we have

$$g_{k+1} - g_k = A(x_{k+1} - x_k) = -\alpha_k A g_k.$$

We can write

$$g_{k+1} = g_k - \alpha_k A g_k$$

- Since A is symmetric, there exists an orthogonal transformation that transforms A into be a diagonal matrix $\Lambda = \text{diag}(\lambda_i)$. Then component-wise,

$$g_i^{k+1} = (1 - \alpha_k \lambda_i) g_i^k, \quad i = 1, 2, \dots, n.$$

We can make two simple observations from the formula

$$g_i^{k+1} = (1 - \alpha_k \lambda_i) g_i^k, \quad i = 1, 2, \dots, n.$$

- ▶ If $g_i^k = 0$ for some k , then all subsequent $g_i^{\tilde{k}} = 0$, $\tilde{k} \geq k$.
- ▶ If $\alpha_k = \lambda_i^{-1}$ at step k , then $g_i^{k+1} = 0$.

Then if we choose $\alpha_i = \lambda_i^{-1}$, $i = 1, 2, \dots, n$, we will get $g_n = 0$.
So let us assume that we cannot somehow get the eigenvalues of A all at once.

- ▶ Now let us introduce the limited-memory setting: suppose n is larger and we store our data in a long vector of size $1 \leq m \leq n$
- ▶ Repeatedly applying the formulae

$$x_k = x_{k-1} - \alpha_{k-1}g_{k-1}, \quad g_k = g_{k-1} - \alpha_{k-1}Ag_{k-1},$$

it is easy to see

$$x_k - x_{k-m} \in \text{span} \{g_{k-m}, Ag_{k-m}, A^2g_{k-m}, \dots, A^{m-1}g_{k-m}\}$$

and

$$g_k - g_{k-m} \in \text{span} \{Ag_{k-m}, A^2g_{k-m}, \dots, A^mg_{k-m}\}.$$

- The columns of

$$K_m = [g_{k-m}, Ag_{k-m}, A^2g_{k-m}, \dots, A^{m-1}g_{k-m}]$$

is in general not orthogonal, but since A is symmetric, we can apply the Lanczos algorithm to find an orthonormal basis of K_m .

► The Lanczos algorithm:

Input:

1. $n \times n$ symmetric matrix A
2. some number m , $1 \leq m \leq n$
3. a starting vector q_1

Output:

1. $n \times m$ matrix Q whose columns q_1, \dots, q_m are orthogonal
2. $m \times m$ matrix $T = Q^T A Q$ whose eigenvalues we will need

The Lanczos algorithm

- ▶ 1. First, normalize q_1 if it is not a unit vector already
- 2. Let $w_1 = Aq_1 - q_1^T Aq_1$
- 3. for $j = 2, \dots, m$:
 - 3.1 Let $\beta_j = \|w_{j-1}\|$
 - 3.2 Let $q_j = w_{j-1}/\beta_j$
 - 3.3 Let $w_j = Aq_j - q_j^T Aq_j - \beta_j q_{j-1}$

- ▶ The eigenvalues of

$$T = Q^T A Q$$

are known as the *Ritz values*. Since T is $m \times m$, there are m such Ritz values.

- ▶ The eigenvalues of

$$T = Q^T A Q$$

are known as the *Ritz values*. Since T is $m \times m$, there are m such Ritz values.

- ▶ It is easy to see that when $m = 1$, the only Ritz value is

$$\theta = \frac{p_{k-1}^T A p_{k-1}}{p_{k-1}^T p_{k-1}}$$

so $1/\theta$ is the Barzilai-Borwein step size.

- ▶ Also, if $m = n$, then the Ritz values are just the eigenvalues of A (this is a property of the Krylov subspace)

The Limited Memory Steepest Descent Method

- ▶ The sequence of steepest descent iterations is divided into groups of m iterations, referred to as *sweeps*. Suppose m Ritz values are calculated

$$\theta_{j,k-1}, \quad j = 1, 2, \dots, m$$

- ▶ At each sweep, we use

$$x^{j+1,k} = x^{j,k} - \alpha_{j,k} g^{j,k}, \quad j = 1, 2, \dots, m$$

where

$$x^{1,k} := x^k, \quad \alpha_{j,k} = (\alpha_{j,k-1})^{-1}, \quad g^{j,k} = Ax^{j,k},$$

and let

$$x^{m+1,k} =: x^{k+1}.$$

Summary of the Algorithm

- ▶ 1. Fix m such that $1 \leq m \leq n$
- 2. Loop the following:
 - 2.1 Use x to define $g = Ax$
 - 2.2 Use (A, m, g) to find m Ritz values θ_j , $j = 1, \dots, m$
 - 2.3 For $j = 1, \dots, m$, do steepest descent with step size $\alpha_j = 1/\theta_j$

Key Ingredients

Recall that we had

$$x_k - x_{k-m} \in \text{span} \{g_{k-m}, Ag_{k-m}, A^2g_{k-m}, \dots, A^{m-1}g_{k-m}\}.$$

- ▶ When $m = 1$, we have the usual steepest descent with the Barzilai-Borwein step size
- ▶ When $m = n$, we have a complete set of eigenvalues for A
- ▶ So we are, in a sense, interpolating between the two cases
- ▶ Also, with the knowledge of m Ritz values, we can jump m iterations each time (from x_{k-m} to x_k), which the author calls a “sweep”

HOME PLOTS APPS EDITOR PUBLISH VIEW

Editor - /Users/yan/MATLAB/sweep_method_quadratic.m

```

1 function [x, fun_val, X] = sweep_method_quadratic(A, c, x0, epsilon)
2 % use steepest descent method to find the optimal solution for following
3 % quadratic model
4 % min f(x) = 0.5 x'*A*x - c'*x
5 % input : A ----- the positive definite matrix associated with the objective
6 % function
7 % c ----- a column vector associated with the linear part of the
8 % objective function
9 % x0 ----- starting point of the method
10 % epsilon ----- tolerance parameter
11
12 % output : x ----- an optimal solution (up to a tolerance)
13 % fun_val --- the optimal value up to a tolerance
14 % X ----- contain all x_k
15
16 X = [];
17 X = [X x0];
18 x = x0;
19 iter = 0;
20 grad = A*x - c;
21 p = -grad;
22 %p = A \ (-grad);
23 fun_val = 0.5*x'*A*x - c'*x;
24 for i = 1:3
25     iter = iter + 1;
26     %alpha = -(grad)'*p/(p'*A*p); % exact line search
27     %x = x + alpha*p;
28     grad = A*x - c;
29     p = -grad;
30
31     q1 = grad / norm(grad);
32     [Q, H] = qr(q1, 'upper');
33     theta = eig(Q' * A * Q);
34
35     for i = 1:length(theta)
36         p = -A*x;
37         x = x + (1 / theta(i)) * p;
38     end
39
40     %p = A \ (-grad);
41     fun_val = 0.5*x'*A*x - c'*x;
42     fprintf('sweep = %3d norm_grad = %2.6f fun_val = %2.6f\n', iter, norm(grad),
43           X = [X x];
44 end

```

Command Window

```

>> test_steepest_descent
0 =
    1     0     0     0
    0     2     0     0
    0     0     3     0
    0     0     0     4

sweep = 1 norm_grad = 5.477226 fun_val = 0.012632
sweep = 2 norm_grad = 0.209852 fun_val = 0.000032
sweep = 3 norm_grad = 0.013837 fun_val = 0.000000

```

UTF-8 sweep_method_quadratic Ln 42 Col 15

Thank you!